



Weekly Progress Report

Project: TinyML-Based Project on FPGA Board with RISC-V Core

Name: Aman Chauhan

Branch: Btech. Computer Science and Engineering (CSE Core)

Roll Number: 22BCE0476

College: Vellore Institute of Technology, Vellore

Mentor: Dr. Sudip Roy

Week: 4th report

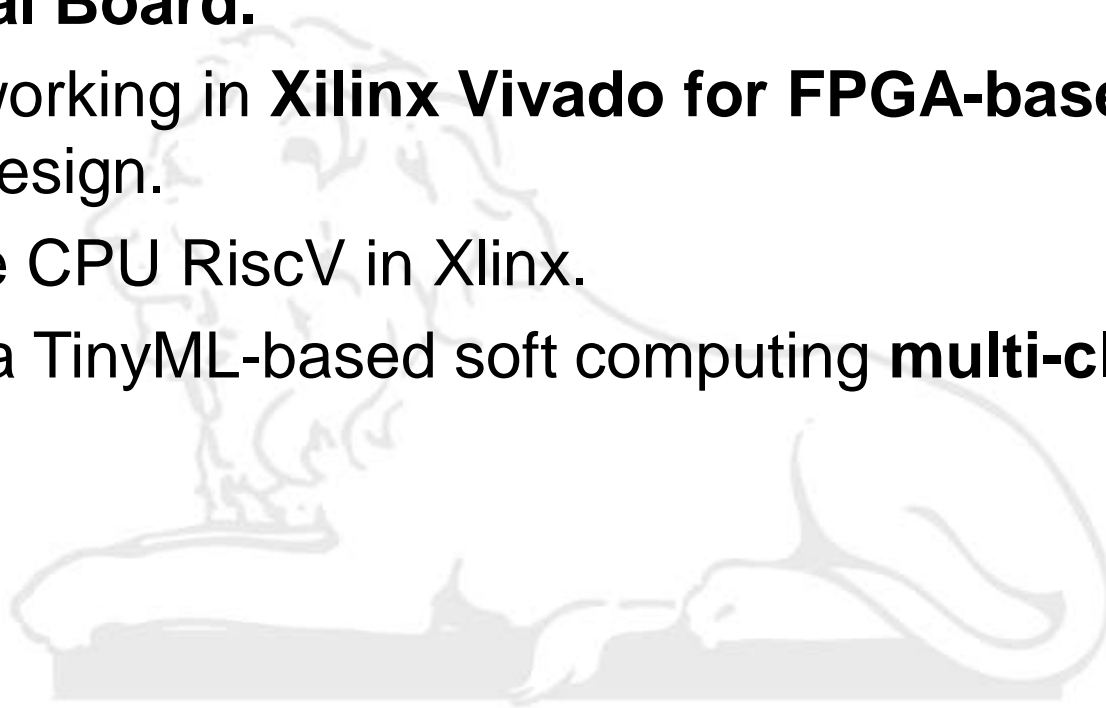
Date: 13th May 2025





Summary of Week 4 Activities

- Completed “**Mixing C assembly functions**” lesson from the Harvard edX RVfpga Course using SweRVolf on the **Nexys A7 Virtual Board**.
- Started working in **Xilinx Vivado for FPGA-based RISC-V** system design.
- Soft Core CPU RiscV in Xlinx.
- Initiated a TinyML-based soft computing **multi-classification model**.



Rvfpga Image Processing



Program that processes an RGB image (left side of the image below), and generates a grayscale version of that image (right side of the image below).



Transformation of an RGB Image to a Grayscale Image

```

49 int main(void) {
50     // Create an NxM matrix using the input image
51     initColourImage(ColourImage);
52
53     // Transform Colour Image to Grey Image
54     ColourToGrey(ColourImage, GreyImage);
55
56     // Initialize Uart
57     uartInit();
58     // Print message on the serial output
59     printfNexys("Created Grey Image");
60
61     while(1);
62
63     return 0;
64 }

```

```

typedef struct {
    unsigned char R;
    unsigned char G;
    unsigned char B;
} RGB;

```

```
unsigned char GreyImage[N][M];
```

2-Dimensional Array of Characters

```

38 extern int ColourToGrey_Pixel(int R, int G, int B);
39
40 void ColourToGrey(RGB Colour[N][M], unsigned char Grey[N][M]) {
41     int i, j;
42
43     for (i=0; i<N; i++)
44         for (j=0; j<M; j++)
45             Grey[i][j] = ColourToGrey_Pixel(Colour[i][j].R, Colour[i][j].G, Colour[i][j].B);
46 }

```

ColourToGrey Function and ColourToGrey_Pixel Subroutine

```

1  .globl ColourToGrey_Pixel
2
3  .text
4
5  ColourToGrey_Pixel:
6
7      li x28, 306
8      mul a0, a0, x28
9
10     li x28, 601
11     mul a1, a1, x28
12
13     li x28, 117
14     mul a2, a2, x28
15
16     add a0, a0, a1
17     add a0, a0, a2
18
19     srl a0, a0, 10
20
21     ret
22
23 .end

```



- **SweRV EH1** lacks floating-point support, so **RGB to Grayscale** conversion uses integer arithmetic.
- Standard weights (0.299, 0.587, 0.114) are scaled by 1024, giving integers: 306 (R), 601 (G), 117 (B).
- Final formula: **Grayscale = (306×R + 601×G + 117×B) >> 10** (equivalent to dividing by 1024).
- Ensures grayscale output stays within 0–255 range.
- Implemented using a C function (ColourToGrey) and an assembly subroutine (ColourToGrey_Pixel).



Xilinx Vivado for FPGA-based RISC-V

- Used **Xilinx Vivado Design Suite** to develop a **RISC-V based system** on an **industrial-grade FPGA board**.
- Created a **custom block design** using Vivado IP Integrator, integrating RISC-V core with AXI interconnects.
- Configured **memory controllers and peripheral interfaces** suitable for SoC-style architecture.
- Synthesized and implemented the design, followed by **bitstream generation** and **hardware debugging**.
- Focused on real-world deployment, gaining insights into **timing analysis, resource utilization, and hardware constraints**.



Vivado 2024.2

Implementation Complete

Default Layout

Flow Navigator

- SIMULATION
 - Run Simulation
- RTL ANALYSIS
- SYNTHESIS
 - Run Synthesis
 - Open Synthesized Design
 - Constraints Wizard
 - Edit Timing Constraints
 - Set Up Debug
 - Open Dataflow Design
 - Report Timing Summary
 - Report Clock Networks
 - Report Clock Interaction
 - Report Methodology
 - Report DRC
 - Report Noise
 - Report Utilization
 - Report Power
 - Schematic

- IMPLEMENTATION
- Run Implementation
- Open Implemented Design
 - Constraints Wizard
 - Open Dataflow Design
 - Edit Timing Constraints
 - Report Timing Summary

Netlist

counter

- Nets (25)
- Leaf Cells (21)

Properties

Select an object to see properties

Project Summary

Device

adder.x

adder_tlv

counter.x

Timing

Design Timing Summary

General Information

Timer Settings

Design Timing Summary

Methodology Summary (4)

Check Timing (23)

Intra-Clock Paths

Inter-Clock Paths

Other Path Groups

Timer Ignored Paths

Timing Summary - impl_1 (saved)

Setup	Hold	Pulse Width
Worst Negative Slack (WNS): inf	Worst Hold Slack (WHS): inf	Worst Pulse Width Slack (WPWS): NA
Total Negative Slack (TNS): 0.000 ns	Total Hold Slack (THS): 0.000 ns	Total Pulse Width Negative Slack (TPWS): NA
Number of Failing Endpoints: 0	Number of Failing Endpoints: 0	Number of Failing Endpoints: NA
Total Number of Endpoints: 15	Total Number of Endpoints: 15	Total Number of Endpoints: NA

There are no user specified timing constraints.



1650365053

vivado_simulation - [C:\Users\Aditya\Documents\vivado_simulation\vivado_simulation.xpf] - Vivado 2024.2

File Edit View Tools Reports Windows Layout View Help

Flow Navigator

- SIMULATION
 - Run Simulation
- RTL ANALYSIS
- SYNTHESIS
 - Run Synthesis
 - Open Synthesized Design
 - Constraints Wizard
 - Edit Timing Constraints
 - Set Up Debug
 - Open Database Design
 - Report Timing Summary
 - Report Clock Networks
 - Report Clock Interaction
 - Report Methodology
 - Report DRC
 - Report Notes
 - Report Utilization
 - Report Power
 - Schematic

- IMPLEMENTATION
- Run Implementation
- Open Implemented Design
 - Constraints Wizard
 - Open Database Design
 - Edit Timing Constraints
 - Report Timing Summary
 - Report Clock Networks
 - Report Clock Interaction
 - Report Methodology

SYNTHESIZED DESIGN - xcp7ay/Rct156-2.e

Netlist

counter

- 0: Nto (20)
- 1: Leaf Cells (21)

Properties

Default all objects to use properties

Project Summary

Device: xcp7ay/Rct156-2.e

adder_0x0

counter_0

Schematic

20 Cells, 8 I/O Ports, 25 Nets

Tcl Console

Messages Log Reports Design Tools Timing

INFO: (Script 29-20) Device Transformation completed in 0.029 seconds
INFO: (Project 1-474) Netlist was created with Vivado 2024.2
INFO: (Project 1-570) Preparing netlist for logic optimization.
INFO: (Step 11-138) Running 0 iterations to 0.1000000000.
Netlist sorting complete. Time (s): cpu = 00:00:00 : elapsed = 00:00:06 : Memory (MB): peak = 5677.067 : gain = 0.508
INFO: (Project 1-111) Device Transformation Summary:
A total of 4 instances were transformed.
INFO => 2025 (ISUPCTRL, D8007): 4 instances

Type = Tcl command line

Netlist: xcp7ay/Rct156-2.e

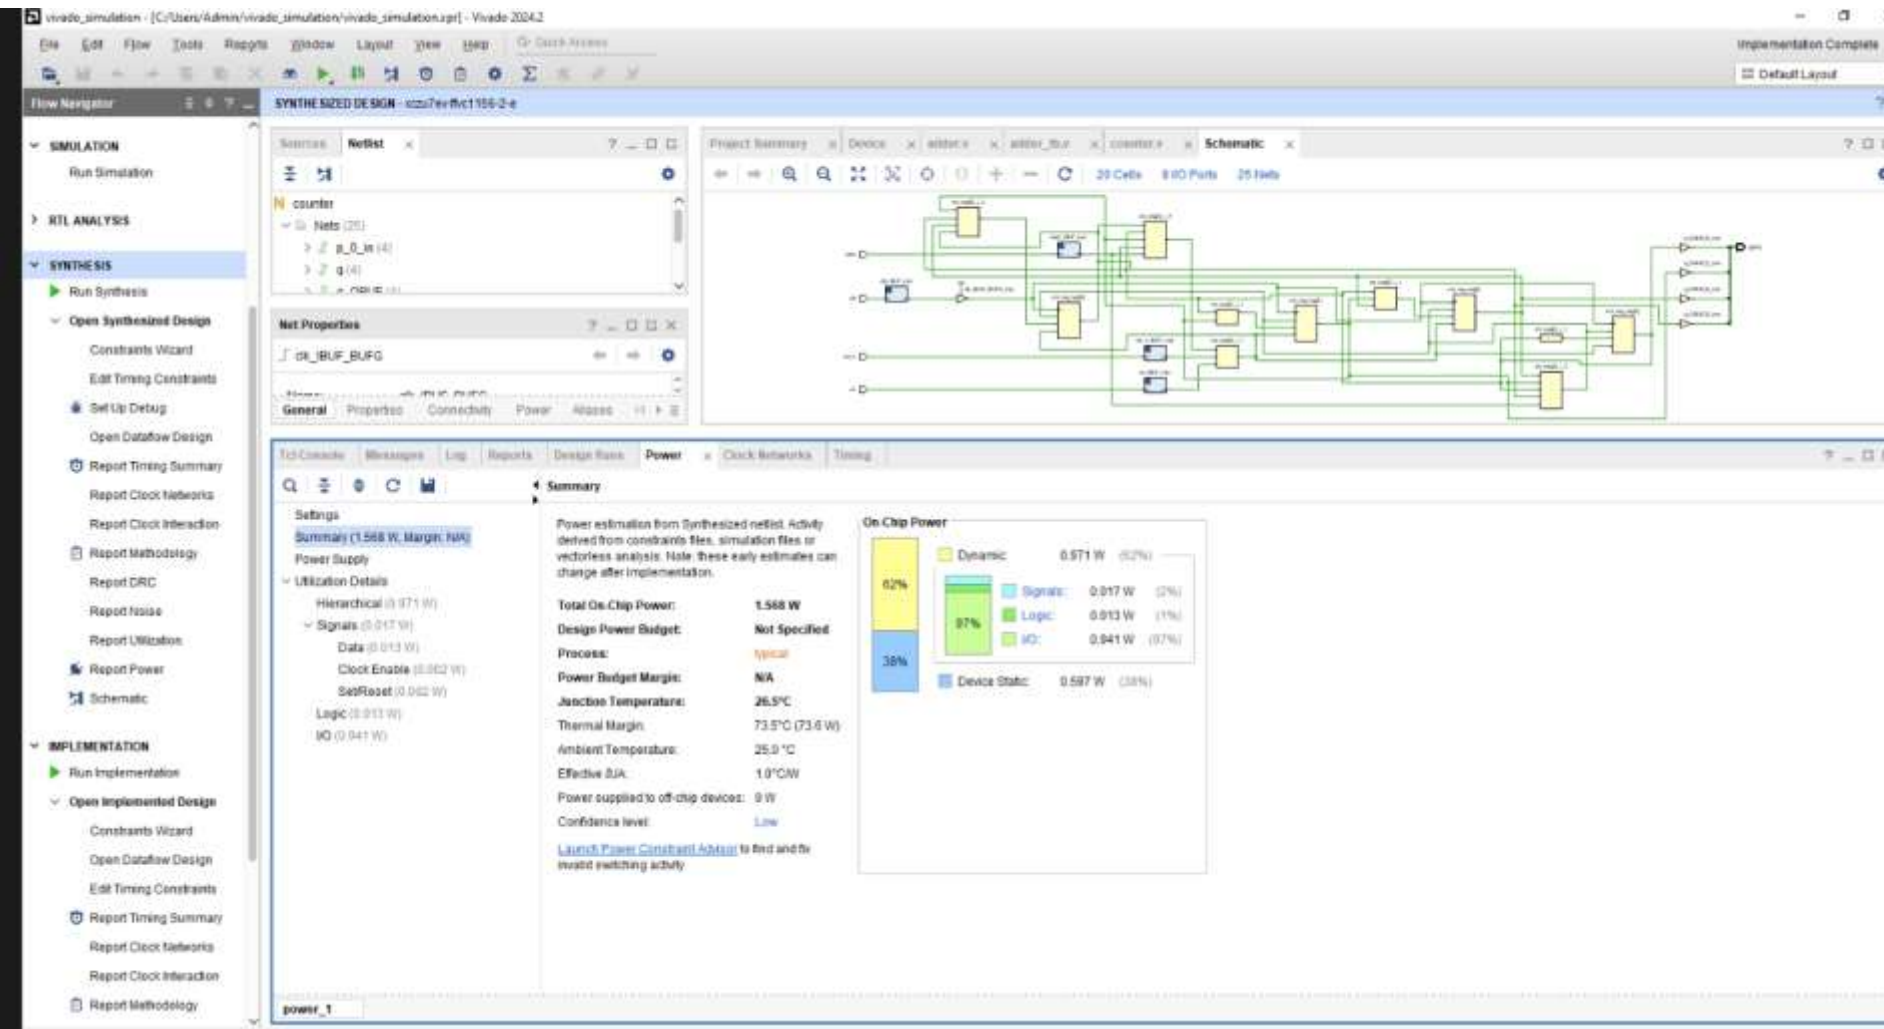
Type: GLOBAL_CLOCK (clock) Route status: Has unplaced ports or pins

Type here to search

OpenAI's Sam Altma...

1760

11-08-2024



Multi-Classification Model



Steps Involved:-

Libraries Installation:-

Oceans cover two-thirds of the planet. In this assignment, you will build a classifier to tell several types of creatures apart.

```
import os

from collections import Counter

import matplotlib.pyplot as plt
import torch
import torch.nn as nn
import torch.optim as optim
from PIL import Image
from sklearn.metrics import ConfusionMatrixDisplay, confusion_matrix
from torch.utils.data import DataLoader, random_split
from torchinfo import summary
from torchvision import datasets, transforms
from tqdm.notebook import tqdm

torch.backends.cudnn.deterministic = True
```

Environment Setup:-

```
if torch.cuda.is_available():  
    device = "cuda"
```

```
os.listdir("sea_creatures")
```

```
['test', 'train']
```

```
train_dir = "sea_creatures/train"
```

```
# Get the list of class names (each folder is a class)
```

```
classes = os.listdir(train_dir)
```

```
# Print the class names
```

```
print(classes)
```

```
['Puffers', 'Sea Urchins', 'Turtle_Tortoise', 'Whale', 'Jelly Fish', 'Sharks', 'Octopus', 'Sea Rays', 'Dolphin']
```

Transform Pipeline:-

```
height = 224  
width = 224  
class ConvertToRGB:  
    def __call__(self, img):  
        if img.mode != "RGB":  
            img = img.convert("RGB")  
        return img  
transform = transforms.Compose([  
    ConvertToRGB(),  
    transforms.Resize((224, 224)),  
    transforms.ToTensor()  
)  
print(transform)
```

```
Compose(  
  <__main__.ConvertToRGB object at 0x7fa102cbbb10>  
  Resize(size=(224, 224), interpolation=bilinear, max_size=None, antialias=True)  
  ToTensor()  
)
```



```
sample_file = "sea_creatures/train/Dolphin/10004986625_0f786ab86b_b.jpg"

image = Image.open(sample_file)

transformed_image = transform(image)
print(transformed_image.shape)

torch.Size([3, 224, 224])

dataset = ImageFolder("sea_creatures/train", transform=transform)
print("Image size", dataset[0][0].shape)
print("Label", dataset[0][1])

Image size torch.Size([3, 224, 224])
Label 0
```

Data loader(batch size-30)

```
batch_size = 32
dataset_loader = DataLoader(dataset, batch_size=batch_size)
# Get one batch
first_batch = next(iter(dataset_loader))
print(f"Shape of one batch: {first_batch[0].shape}")
print(f"Shape of labels: {first_batch[1].shape}")

Shape of one batch: torch.Size([32, 3, 224, 224])
Shape of labels: torch.Size([32])
```



Transform Normalize:-

```
transform_norm = transforms.Compose([
    ConvertToRGB(),
    transforms.Resize((224, 224)),
    transforms.ToTensor(),
    transforms.Normalize(mean=mean, std=std),
])
print(transform_norm)

Compose(
  <__main__.ConvertToRGB object at 0x7fa100bbda50>
  Resize(size=(224, 224), interpolation=bilinear, max_size=None, antial
  ToTensor()
  Normalize(mean=tensor([0.2992, 0.4125, 0.4588]), std=tensor([0.2697,
  )

norm_dataset = datasets.ImageFolder(root=train_dir, transform=transform_no
print("Image size", norm_dataset[0][0].shape)
print("Label", norm_dataset[0][1])

Image size torch.Size([3, 224, 224])
Label 0
```

Set up data loaders for both the training and validation data sets. Use the same batch size as before. Remember to set `shuffle=True` on the training loader.

```
train_loader = DataLoader(train_dataset, batch_size=32, shuffle=True)

val_loader = DataLoader(val_dataset, batch_size=batch_size)
```



- Build Model:-

```
model.append(torch.nn.Dropout(p=0.5))
model.append(torch.nn.Linear(in_features=576, out_features=500))
model.append(torch.nn.ReLU())
model.append(torch.nn.Dropout())
model.append(torch.nn.Linear(500, 9)) # 9 output classes

summary(model, input_size=(batch_size, 3, height, width))
```

```
=====
Layer (type:depth-idx)                               Output Shape                                Param #
=====
Sequential                                             [32, 9]                                    --
├─Conv2d: 1-1                                           [32, 16, 224, 224]                        448
├─ReLU: 1-2                                             [32, 16, 224, 224]                        --
├─MaxPool2d: 1-3                                       [32, 16, 56, 56]                         --
├─Conv2d: 1-4                                           [32, 32, 56, 56]                         4,640
├─ReLU: 1-5                                             [32, 32, 56, 56]                         --
├─MaxPool2d: 1-6                                       [32, 32, 14, 14]                         --
├─Conv2d: 1-7                                           [32, 64, 14, 14]                         18,496
├─ReLU: 1-8                                             [32, 64, 14, 14]                         --
├─MaxPool2d: 1-9                                       [32, 64, 3, 3]                           --
├─Flatten: 1-10                                        [32, 576]                                 --
├─Dropout: 1-11                                        [32, 576]                                 --
├─Linear: 1-12                                         [32, 500]                                 288,500
├─ReLU: 1-13                                           [32, 500]                                 --
├─Dropout: 1-14                                        [32, 500]                                 --
├─Linear: 1-15                                         [32, 9]                                   4,509
=====
Total params: 316,593
Trainable params: 316,593
Non-trainable params: 0
```




Cross Entropy Loss:-

```
loss_fn = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=0.001)
model.to(device)
# Send the model to the GPU
```

```
Sequential(
  (0): Conv2d(3, 16, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (1): ReLU()
  (2): MaxPool2d(kernel_size=4, stride=4, padding=0, dilation=1, ceil_mode=False)
  (3): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (4): ReLU()
  (5): MaxPool2d(kernel_size=4, stride=4, padding=0, dilation=1, ceil_mode=False)
  (6): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (7): ReLU()
  (8): MaxPool2d(kernel_size=4, stride=4, padding=0, dilation=1, ceil_mode=False)
  (9): Flatten(start_dim=1, end_dim=-1)
  (10): Dropout(p=0.5, inplace=False)
  (11): Linear(in_features=576, out_features=500, bias=True)
  (12): ReLU()
  (13): Dropout(p=0.5, inplace=False)
  (14): Linear(in_features=500, out_features=9, bias=True)
)
```



Training Of Dataset

```
# Import the train and predict functions from `training.py`, instead of typing them out!
from training import train, predict
epochs = 10
train(model,optimizer,loss_fn,train_loader,val_loader,epochs=10,device=device)
# Train the model for 10 epochs
```

```
Training:  0%|          | 0/155 [00:00<?, ?it/s]
Scoring:   0%|          | 0/39 [00:00<?, ?it/s]
Epoch: 1, Training Loss: 1.69, Validation Loss: 1.50, Validation accuracy = 0.47
Training:  0%|          | 0/155 [00:00<?, ?it/s]
Scoring:   0%|          | 0/39 [00:00<?, ?it/s]
Epoch: 2, Training Loss: 1.46, Validation Loss: 1.40, Validation accuracy = 0.51
Training:  0%|          | 0/155 [00:00<?, ?it/s]
Scoring:   0%|          | 0/39 [00:00<?, ?it/s]
```

Evaluate Model Performance:-

```
# Compute the probabilities for each validation image
probabilities = predict(model,val_loader,device)
# Get the index associated with the largest probability for each
predictions = torch.argmax(probabilities,dim=1)

print("Number of predictions:", predictions.shape)
```

```
Predicting:  0%|          | 0/39 [00:00<?, ?it/s]
Number of predictions: torch.Size([1236])
```

Testing of dataset:-

```
test_dir = "sea_creatures/test"
test_transforms = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.ToTensor()
])

test_dataset = test_dataset = datasets.ImageFolder(root=test_dir, transform=test_transforms)

print("Number of test images:", len(test_dataset))

test_loader = DataLoader(test_dataset, batch_size=32, shuffle=False)
```

Number of test images: 699

```
# Predict the probabilities for each test image
test_probabilities = predict(model, test_loader, device=device)

# Get the index associated with the largest probability for each test image
test_predictions = torch.argmax(test_probabilities, dim=1)

print("Number of predictions:", test_predictions.shape)

Predicting:  0%|          | 0/22 [00:00<?, ?it/s]
Number of predictions: torch.Size([699])
```

Task 1.5.22: Convert the class index to the class name for each test image.

```
test_classes = [classes[i] for i in test_predictions]

print("Number of class predictions:", len(test_classes))
```

Predictions through the validation dataset

```
import matplotlib.pyplot as plt
import random

# Sample 12 random indices from the test dataset
sample_indices = random.sample(range(len(test_loader.dataset.samples)), 12)

# Create a grid of 4x3 subplots
fig, axes = plt.subplots(4, 3, figsize=(20, 10))

# Iterate over the sampled indices and plot the corresponding images
for ax, idx in zip(axes.flatten(), sample_indices):
    image_path = test_loader.dataset.samples[idx][0]
    img = Image.open(image_path)

    # Display the image on the axis
    ax.imshow(img)
    ax.axis('off')

    # Get the predicted class for this image
    predicted_class = test_classes[idx]

    # Set the title of the subplot to the predicted class
    ax.set_title(f"Predicted: {predicted_class}", fontsize=14)

plt.tight_layout()
```

Predicted: Turtle_Tortoise



Predicted: Turtle_Tortoise

Predicted: Turtle_Tortoise



Predicted: Whale

Predicted: Sea Urchins



Predicted: Turtle_Tortoise



References

- **RVfpga HarvardX edX and RVfpga source code**
- **TinyML foundational courses and TensorFlow Lite Micro documentation.**
- **WorldQuant University** Applied AI Lab content
- **RVfpga simulator Vibodo, Rvfpga Nexys Board, Pipelines, Whisper** (Debugging and simulation) and **Tracer (GTK wave usage)**, then **whisper for c program files** compilation.
- **Report and internal study materials.**
